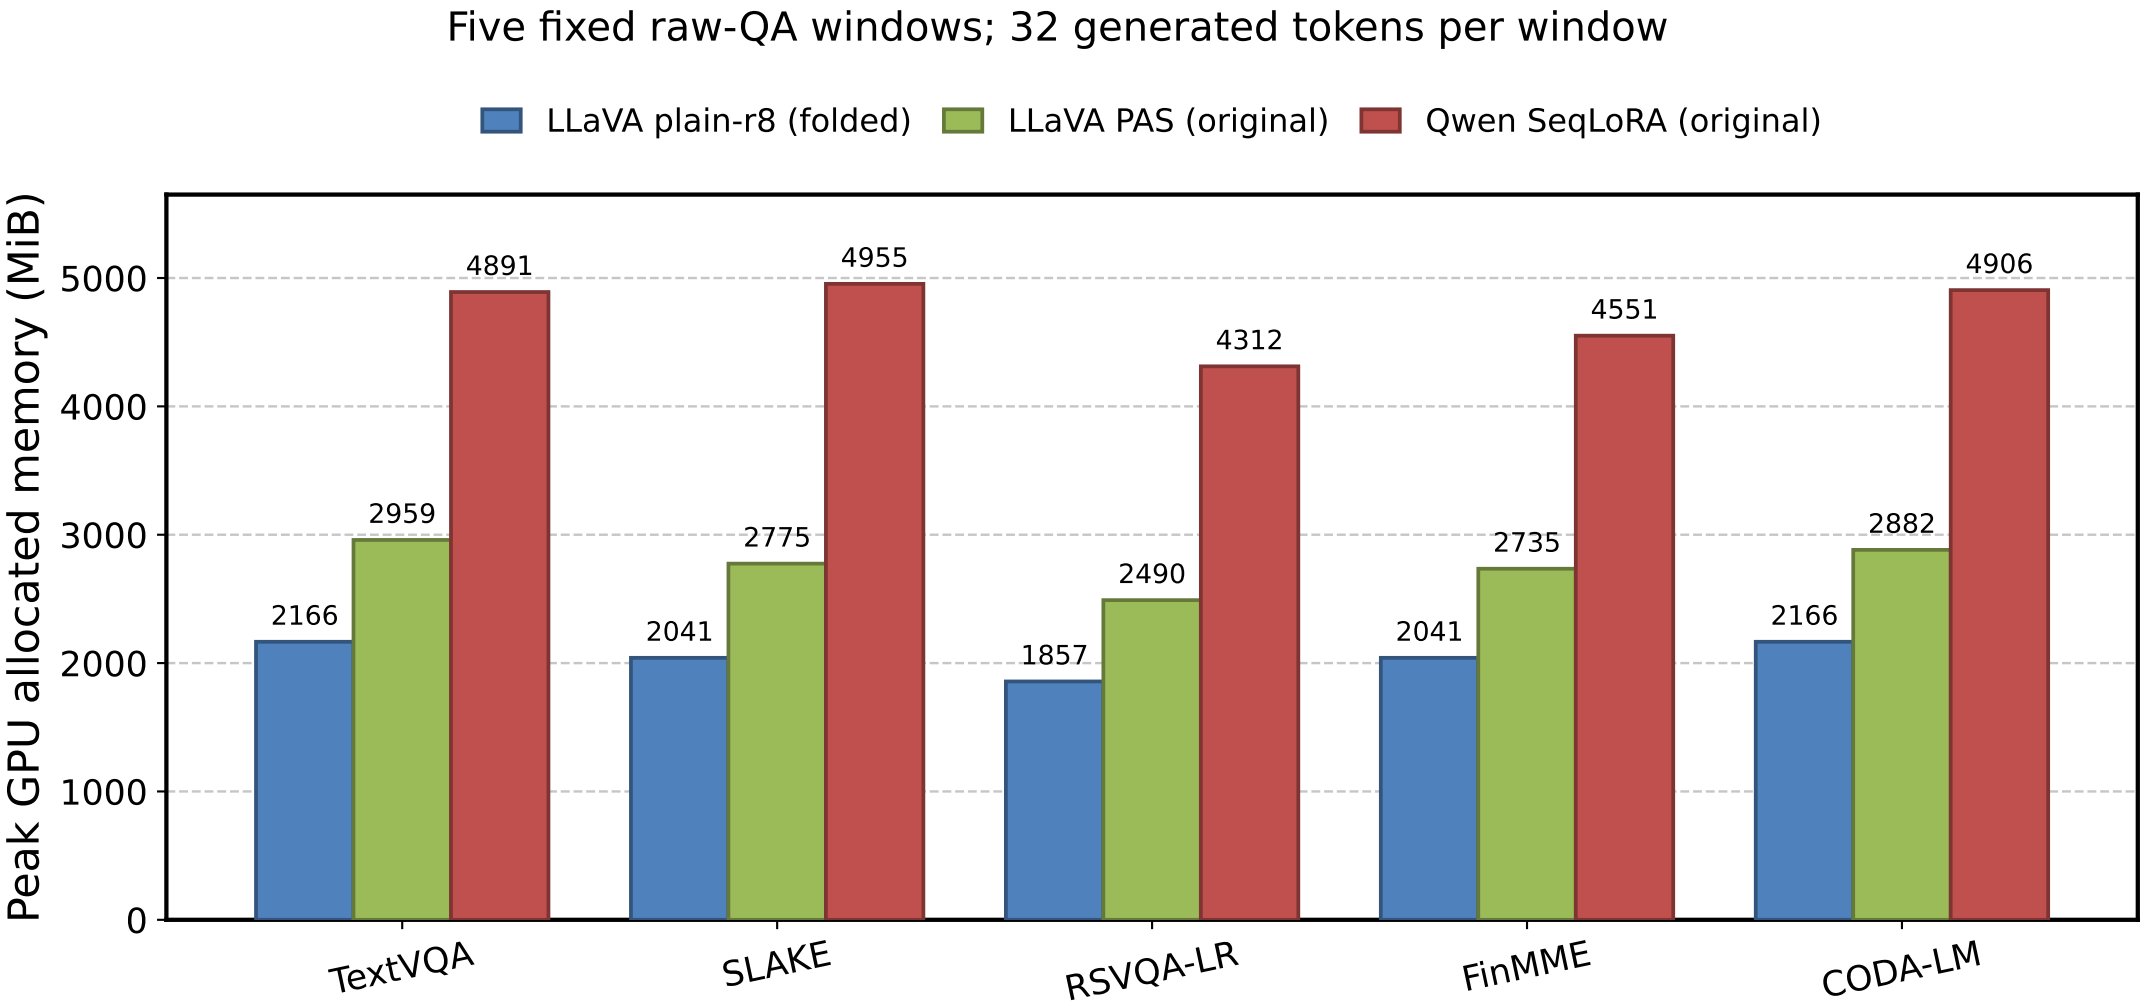

Same physical GPU; native input expansion and original dtypes. Includes input preparation.
Plain-r8: folded FP16. PAS / Qwen: original, unfolded mixed-dtype graphs.
One QA per dataset; no uncertainty estimates. Not an optimal-deployment or full-test peak.